

0.1 Linear Regression

0.1.1 Abstract

Hello, everyone. I'm btobab.

When I learned NLP, I found I completely can't understand formula derivation, so I decided to learn the theoretical derivation of machine learning from the beginning.

This chapter is mainly divided into two parts:

The first part will derive the closed-form solution of the linear regression(least squares method) from three perspectives of matrix, geometry, probability, and provide reference code.

The second part will derive the closed-formula solution of least squares method with regularization from two perspectives of matrix and probability, and construct a complete linear regression class, and implement the closed-solution method and gradient descent method with code at the same time.

0.1.2 Introduction

Linear regression model is a model that use linear function to fit the relationship between one or more independent variables and the dependent variable (y).

The target variable (y) is a continuous numerical type, such as: housing price, number of people and rainfall. The regression model is to find a mapping function between input variables and output variables.

The learning of regression task equals to function fitting: use a function curve to make it fit the known data well and predict unknown data.

Regression task is divided into two processes: model learning and prediction. Construct a model based on given training dataset and predict corresponding output based on new input data.

0.1.3 Algorithm without regularization

Matrix perspective

Note:in general, the vectors we are discussing are column vectors. Therefore, in order to ensure the shape of the matrix during the derivation process, a large number of transposition characters are used

given dataset $\mathcal{D} = \{(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)\}$

in which $x_i \in \mathcal{R}^p, y_i \in \mathcal{R}, i = 1, 2, \dots, n$

$$X = (x_1, x_2, \dots, x_n)^T = \begin{pmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ x_{n1} & x_{n2} & \dots & x_{np} \end{pmatrix}_{np}$$

$$Y = (y_1, y_2, \dots, y_n)^T_{n1}$$

It is the model we construct: $f(w) = w^T x + w_0 x_0$.

Generally set $x_0 = 1$, and $b = w_0 x_0$, b is bias, w is weight. below, for the convenience of derivation, we merge w_0 into w and x_0 into x .

so the model is updated to $f(w) = w^T x$ The loss function of least squares method is:

$$\begin{aligned} L(w) &= \sum_{i=1}^{i=n} \|y_i - w^T x_i\|_2^2 \\ &= (y_1 - w^T x_1 \quad y_2 - w^T x_2 \quad \dots \quad y_n - w^T x_n) \begin{pmatrix} y_1 - w^T x_1 \\ y_2 - w^T x_2 \\ \cdot \\ \cdot \\ y_n - w^T x_n \end{pmatrix} \\ &= (Y^T - w^T X^T)(Y^T - w^T X^T)^T \\ &= (Y^T - w^T X^T)(Y - Xw) \\ &= Y^T Y - w^T X^T Y - Y^T Xw + w^T X^T Xw \end{aligned}$$

Note the second and third terms are transposed to each other, and observe its matrix shape: $(1, p)(p, n)(n, 1) = (1, 1)$

Knowing that these two terms are scalars, and the transposition of a scalar is itself, so the two can be combined, get:

$$L(w) = Y^T Y - 2w^T X^T Y + w^T X^T Xw$$

so $\hat{w} = \text{argmin}(L(w))$ below, to find out the minimum of $L(w)$, we need to differentiate $L(w)$

Note there are three terms in the formula. The first term has nothing to do with w and can be removed. Then the remaining two terms involve matrix derivation.

Regarding to matrix derivation, author recommends three articles by a blogger (more detailed and rigorous than textbook, each formula has proof)

- essence

- basics
- advanced

The following is the derivative solution process of the above two terms.

Because X, Y are constant matrices, the derivative can be obtained directly. However, since it is the derivative of w , the result must be transposed.

$$\frac{d(2w^T X^T Y)}{dw} = 2X^T Y$$

Below, let's solve the third term.

$$\begin{aligned} d(w^T X^T X w) &= tr(d(w^T X^T X w)) = tr(X^T X d(w^T w)) \\ &= tr(X^T X (d(w^T) w + w^T d(w))) = tr(X^T X w (dw)^T) + tr(X^T X w^T dw) \\ &= tr(w^T X^T X dw) + tr(X^T X w^T dw) = tr(2X^T X w^T dw) \end{aligned}$$

so

$$\frac{d(w^T X^T X w)}{dw} = 2w X^T X$$

From a geometric perspective, we regard X as a p dimensional vector.

The first dimension of X is $(x_{11}, x_{21}, \dots, x_{n1})$, the p -th dimension of X is $(x_{1p}, x_{2p}, \dots, x_{np})$

and here Y is regarded as a one-dimensional vector.

$$\text{so } \frac{dL(w)}{dw} = 2X^T X w - 2X^T Y$$

set the derivative equal to 0 to get the closed-form solution of the least squares method:

$$\hat{w} = (X^T X)^{-1} X^T Y$$

Geometry perspective

$$X = (x_1, x_2, \dots, x_n)^T = \begin{pmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ x_{n1} & x_{n2} & \dots & x_{np} \end{pmatrix}_{np}$$

$$Y = (y_1, y_2, \dots, y_n)_{n1}^T$$

Now we assume $p = 2$ because it is easier to draw. The diagram is as follows(I really drew it for a long time and pitifully asked for a star. Change the model to $f(w) = Xw$, which means zoom X with weight w




Figure 1: figure 1

The geometric meaning of the least squares method is to find a w , so that the distance between vector $Y - Xw$ and space w is the smallest. Of course the case of the smallest distance is perpendicular to space X .

so we get a formula: $X^T(Y - Xw) = 0$

then get the solution of w :

$$\begin{aligned} X^T X w &= X^T Y \\ \hat{w} &= (X^T X)^{-1} X^T Y \end{aligned}$$

we can see that the solved w is the same as the result of matrix perspective.

Probability perspective

As we known, in reality, it is hard to fit a distribution with a straight line. True data must have some randomness, that is, noise.

so we assume noise $\epsilon \sim N(0, \sigma^2)$

so $y = f(w) + \epsilon = w^T x + \epsilon$

so $y|x; w \sim N(w^T x, \sigma^2)$

Bring it into the probability density function of gaussian distribution:

$$p(y|x; w) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y - w^T x)^2}{2\sigma^2}}$$

Then use *MLE* (maximum likelihood estimate)

Note: so-called MLE is to get the relative frequency via a large number of samples to approximate probability

Let's assume a function: $\zeta(w) = \log p(Y|X; w)$

Since n data are independent, we can change the probability to a form of continuous multiplication:

$$\zeta(w) = \log \prod_{i=1}^n p(y_i|x_i; w) = \sum_{i=1}^n \log p(y_i|x_i; w)$$

bring the probability density function of gaussian distribution into the formula:

$$\zeta(w) = \sum_{i=1}^n \left(\log \frac{1}{\sqrt{2\pi}\sigma} - \frac{(y - w^T x)^2}{2\sigma^2} \right)$$

since the former term has nothing to do with w , it can be ignored.

so:

$$\begin{aligned} \hat{w} &= \operatorname{argmax} \zeta(w) \\ &= \operatorname{argmax} \sum_{i=1}^n - \frac{(y - w^T x)^2}{2\sigma^2} \\ &= \operatorname{argmin} \sum_{i=1}^n (y - w^T x)^2 \end{aligned}$$

The conclusion obtained by using maximum likelihood estimation is the definition of least squares method.

This also shows that least squares method hides a assumption that noise is gaussian distribution.

0.1.4 Implement without regularization

```
%matplotlib inline
import numpy as np
import matplotlib.pyplot as plt

# num of samples
n = 1000
# noise
epsilon = 1
X = np.expand_dims(np.linspace(0,100,1000), axis=-1)
w = np.asarray([5.2])
Y = X * w
# apply noise to X
X += np.random.normal(scale=epsilon, size=(X.shape))
X_T = X.transpose()
w_hat = np.matmul(np.linalg.pinv((np.matmul(X_T, X))), np.matmul(X_T, Y))
print(w_hat)
plt.scatter(X, Y, s=3, c="y")
Y_hat = X * w_hat
plt.plot(X, Y_hat)
plt.show()
```

0.1.5 Algorithm with regularization

Matrix perspective

Firstly, given a new loss function with regularization:

$$\zeta(w) = \sum_{i=1}^n \|y_i - w^T x_i\|^2 + \lambda \|w\|^2$$

then the derivation of loss function from matrix perspective without regularization is referenced:

$$\zeta(w) = Y^T Y - 2w^T X^T Y + w^T X^T X w + \lambda \|w\|^2$$

so $\hat{w} = \operatorname{argmax}(\zeta(w))$

differentiate $\zeta(w)$:

$$\frac{\partial \zeta(w)}{\partial w} = 2X^T X w - 2X^T Y + 2\lambda w$$

set the derivative equal to 0 to get the closed-form solution of least squares method with regularization:

$$\hat{w} = (X^T X + \lambda I)^{-1} X^T Y$$

I is identity matrix

Probability perspective

assume noise $\epsilon \sim N(0, \sigma_1^2)$ $w \sim N(0, \sigma_2^2)$

since $y = w^T x + \epsilon$

we get $y|w \sim N(w^T x, \sigma_1^2)$

Next we use MAP(Maximum a posteriori estimate):

according to Bayes theorem:

$$P(w|Y) = \frac{P(Y|w)P(w)}{P(Y)}$$

$P(w)$ is a priori probability, $P(Y|w)$ is a likelihood probability, $P(Y)$ is normalized probability, prior probability is multiplied by the likelihood probability and normalized to obtain the posterior probability $P(w|Y)$. actually $P(Y)$ is constant, so:

$$\hat{w} = \operatorname{argmax}(P(w|Y)) = \operatorname{argmax}(P(Y|w)P(w)) = \operatorname{argmax}(\log(P(Y|w)P(w)))$$

since samples are independent, we can change the probability to a form of continuous multiplication.

$$= \operatorname{argmax}(\log(\prod_{i=1}^n P(y_i|w)P(w))) = \operatorname{argmax}(\sum_{i=1}^n \log(P(y_i|w) + \log(P(w))))$$

bring it into probability density function of gaussian distribution to get:

$$\hat{w} = \operatorname{argmax}(\sum_{i=1}^n \log(\frac{1}{\sqrt{2\pi}\sigma_1}) - \frac{(y_i - w^T x_i)^2}{2\sigma_1^2} + \log(\frac{1}{\sqrt{2\pi}\sigma_2}) - \frac{w^2}{2\sigma_2^2})$$

since both σ_1 and σ_2 are hyperparameters, they can be omitted.

so:

$$\begin{aligned}\hat{w} &= \operatorname{argmin}(\sum_{i=1}^n \frac{(y_i - w^T x_i)^2}{2\sigma_1^2} + \frac{w^2}{2\sigma_2^2}) \\ &= \operatorname{argmin}(\sum_{i=1}^n (y_i - w^T x_i)^2 + \frac{\sigma_1^2}{\sigma_2^2} w^2)\end{aligned}$$

we can see that the result derived via MAP is the definition of the least squares method with regularization.

0.1.6 Implement with regularization

```
import os
os.chdir("../")
from models.linear_models import LinearRegression
import numpy as np
import matplotlib.pyplot as plt
X_ = np.expand_dims(np.linspace(0, 10, 1000), axis=-1)
X = np.c_[X_, np.ones(1000)]
w = np.asarray([5.2, 1])
Y = X.dot(w)
X = np.r_[X, np.asarray([[11, 1], [12, 1], [13, 1]])]
Y = np.r_[Y, np.asarray([100, 110, 120])]

model = LinearRegression(l2_ratio=1e1, epoch_num=1000, lr=1e-2, batch_size=100,
model.fit(X[:, :-1], Y)
print(model.get_params())
model.draw(X[:, :-1], Y)
```